



Big cats classifier

Dataset

kaggle

10 Big Cats of the Wild -
Image Classification



Used Technologies

The logo for Google Colab, featuring the word "colab" in a bold, lowercase, orange sans-serif font.

colab



Learning parameters


Parameters:

- Batch size
- Epochs
- Learning rate
- Criterion (CrossEntropyLoss)

Optimizers:

- Adam
- SGD (Stochastic Gradient Descent)





Data preparation



Training loop



```
def train(log_file):
    for epoch in range(num_epochs):
        model.train()

        for batch_idx, (inputs, targets) in enumerate(train_loader):
            inputs, targets = inputs.to(device), targets.to(device)
            optimizer.zero_grad()           # zero gradient
            outputs = model(inputs)          # Invoke forward function
            loss = criterion(outputs, targets) # Calculate loss function
            loss.backward()                  # Backpropagation
            optimizer.step()                 # Update weights

        log_file.write(f'{epoch+1},{loss.item():.4f},')
        test_accuracy(log_file)
    if (epoch % 10 == 0):
        print(epoch)
```

Neural network architecture

First network

```
class BigCatClassifier(nn.Module):
    def __init__(self):
        super(BigCatClassifier, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1)
        self.relu1 = nn.ReLU()
        self.maxpool1 = nn.MaxPool2d(kernel_size=2, stride=2)


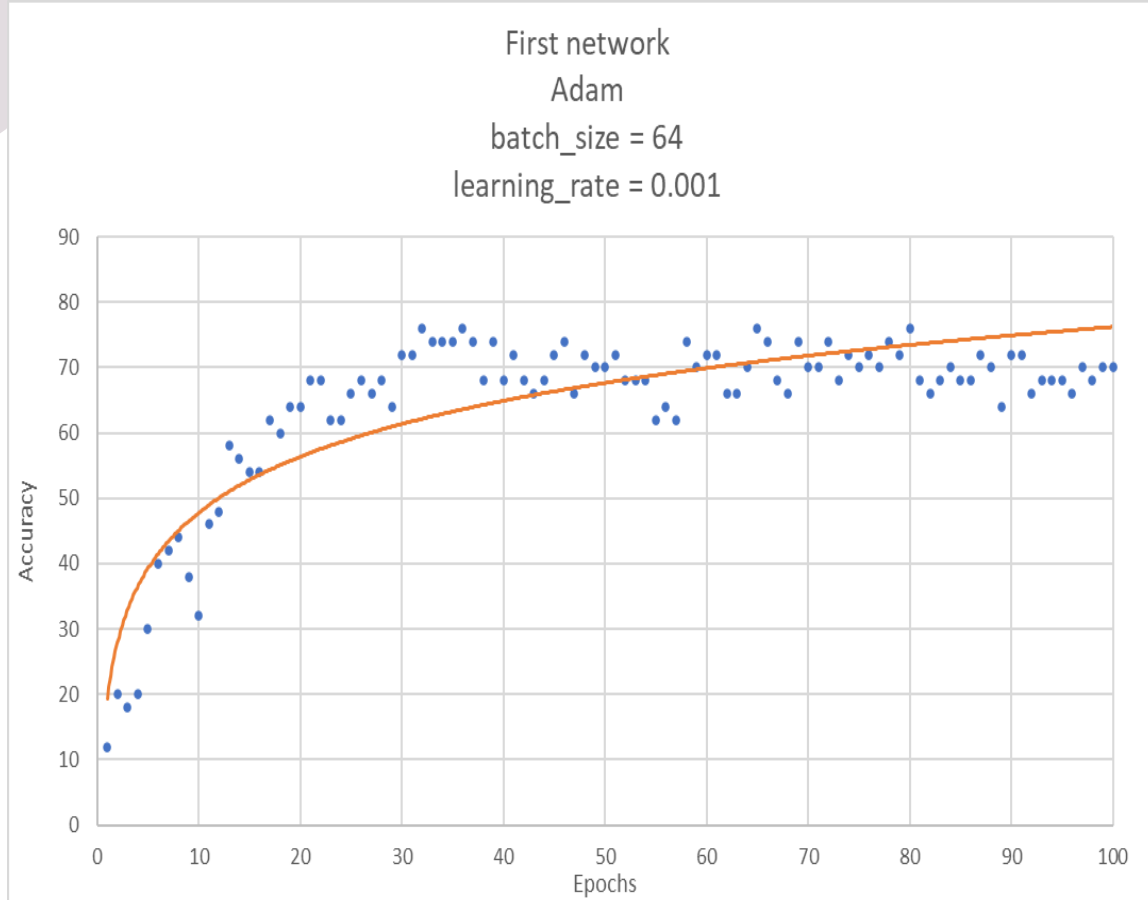
        self.conv2 = nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1)
        self.relu2 = nn.ReLU()
        self.maxpool2 = nn.MaxPool2d(kernel_size=2, stride=2)

        self.flatten = nn.Flatten()

        self.fc1 = nn.Linear(32*56*56, 256)
        self.relu3 = nn.ReLU()
        self.dropout = nn.Dropout(0.5)

        self.fc2 = nn.Linear(256, 10)
```

First net results



Neural network architecture

Second network

```
class BigCatClassifier2(nn.Module):
    def __init__(self):
        super(BigCatClassifier2, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1)
        self.relu1 = nn.ReLU()
        self.maxpool1 = nn.MaxPool2d(kernel_size=2, stride=2)

        self.conv2 = nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1)
        self.relu2 = nn.ReLU()
        self.maxpool2 = nn.MaxPool2d(kernel_size=2, stride=2)



        # Additional convolution layer
        self.conv3 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.relu3 = nn.ReLU()
        self.maxpool3 = nn.MaxPool2d(kernel_size=2, stride=2)

        self.flatten = nn.Flatten()

        self.fc1 = nn.Linear(64*28*28, 256)
        self.relu5 = nn.ReLU()
        self.dropout = nn.Dropout(0.5)

        self.fc2 = nn.Linear(256, 10)
```

Second net results



Next steps



Thanks for listening

